

---

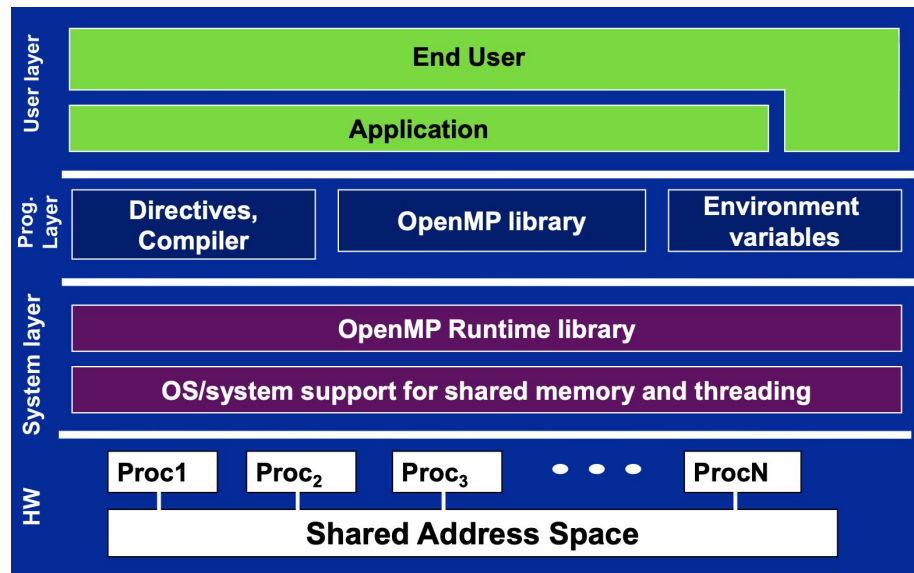
# OMP Programming

Computhon2021-1  
By: Amro Alabsi Aljundi

---


# Introduction - 1

- It is a **parallel programming model** that is very frequently used today.
  - Alternatives to C++ threads or Pthreads.
- **Easy** API composed of **compiler directive** and **library functions**.




# Introduction - 2

- Based on the **shared memory multiprocessing (SMP)** where threads share an address space



**SMP**



**Message Passing  
Interface**

# API Overview


- The API is mostly composed of **compiler directives**.
- These directives act on **structured blocks** (code within { }) that has a single entry and a single exit location

Example:

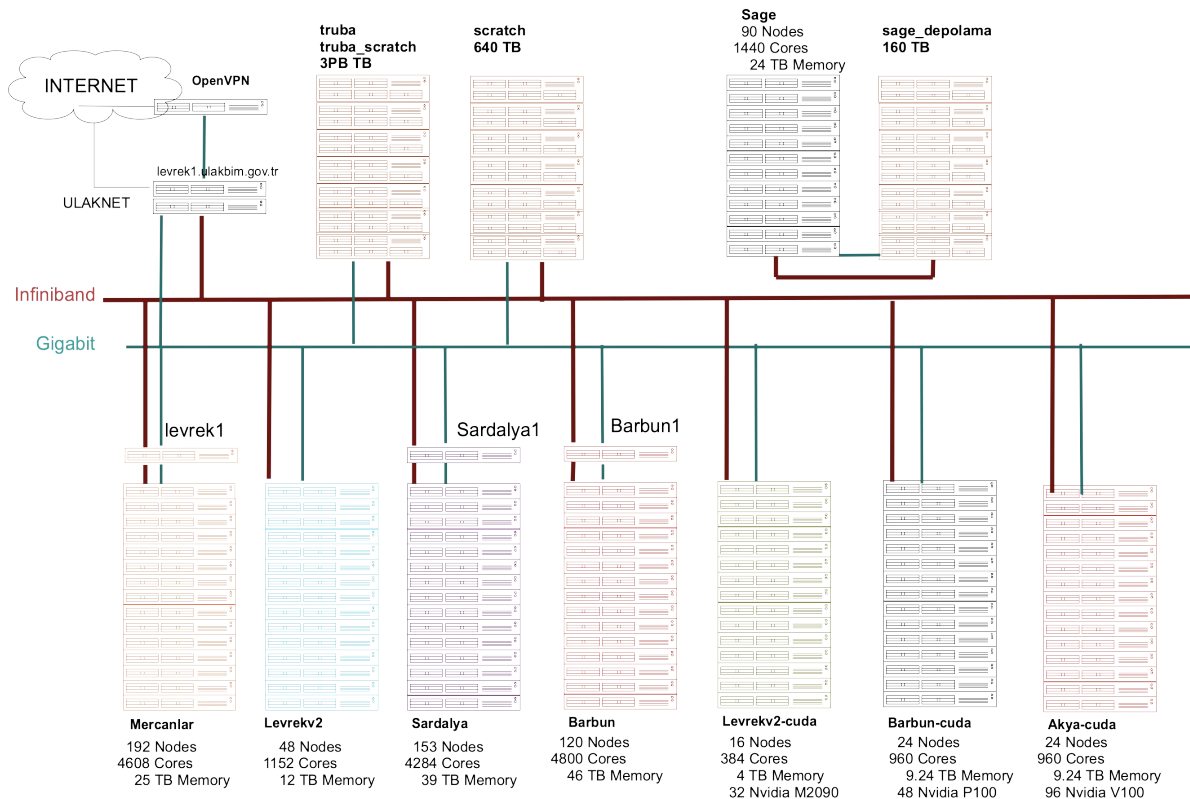
```
#pragma omp parallel num_threads(16)
{
    // entry
    ...
    ...
    // exit
}
```

# Execution model: fork-join

- Execution starts with a single **master thread**.
- It will **fork a team of threads** that run in parallel.
- Team of threads will **join** the master thread.



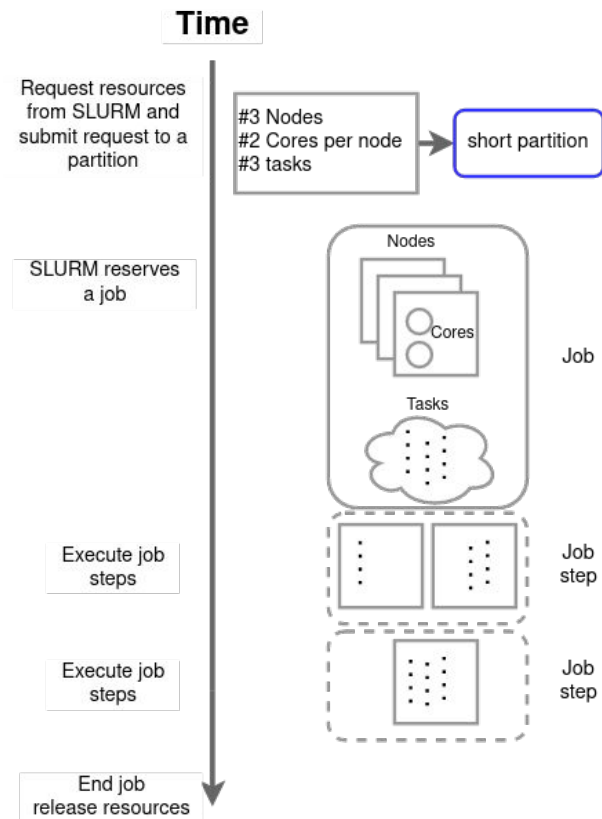
# TRUBA Architecture



# SLURM Model

1. Request a bunch of resources from SLURM
2. SLURM allocates resources by creating a **job** for a specific **time period**.
  - a. Nodes (servers)
  - b. Cores
  - c. Memory
  - d. Tasks (program executions - sort of like processes)
3. Execute work using a job by dispatching **job steps**: program **executions** that use your reserved **tasks**.

**Note: the default allocation time period is 2 minutes.**



# Example 1: creating a parallel region

- `#pragma omp parallel` directive
- Setting number of threads
  - Environment variables
  - `omp_set_num_threads()`
  - `num_threads()` clause of `#pragma omp`
- Nested parallelism




# Running example: calculating Pi

Calculating the following integral returns the value of Pi

$$\int_0^1 \frac{4.0}{(1+x^2)} dx = \pi$$

We can calculate this integral numerically very easily:



# Attempt 0: naive parallelization

- Parallelization of the loop
- Function calls
  - `omp_num_threads()`
  - `omp_thread_num()`
- Variable scopes

# Attempt 1: padding to solve false sharing

- Padding the array solves our issue

## Attempt 2: use a lock to control a single shared variable

- We can use a single shared variable and control its access with a lock
- Where we place the lock is very important!!

## Attempt 3: use private variables for per-thread sums

- There is no reason to access the shared variable at every iteration.
- Each thread can use a variable in its private memory
- Access to the shared variable can be protected with
  - Lock
  - Critical section
  - Atomic

## Attempt 4: use a `#pragma omp for`

- We can leave the task of carrying out the loop to OMP
- We can control the schedule

## Attempt 5: use a `#pragma omp for with reduction`

- We can also ask OMP to carry out the reduction for us.

# Conclusion

- OMP is a parallel programming model for shared-memory systems
- Easy to use API composed of (mainly) compiler directives
- Uses the fork-join execution model

## **Pi parallelization example:**

- Using shared values and false sharing
- Control access to shared variables with locks, critical sections, and atomic
- Worksharing construct: for loop
- Reduction in the for loop